

#llm#agents#interpretability#evaluation#open-source#security#benchmarks#open-weights#rag#reasoning#tool-use#vision

Top tags on this page — every tag on a card is clickable.

AI agents rapidly exploit security vulnerabilities from patch discussions

Modern AI coding agents can now discover and exploit a vulnerability from the public discussion around its patch, well before the fix has reached most deployments.

- Patch discussions leak exploitable detail before rollout.
- The gap between disclosure and a working exploit is closing.
- Open-source projects carry the exposure first.

#security#ai-agents#vulnerability#open-source#exploit

NEWS

relevance 9/10 · Simon Willison · 2026-08-28

[simonwillison.net/2026/Aug/28/just-a-rumour-of-a-bug/ ↗](#)

Sparse autoencoder features transfer unevenly across model scale

A study of how features found by a sparse autoencoder on a small research model correspond to features in a deployed model of the same family, and where that correspondence breaks down.

- The clearest correspondences are the narrowest features.
- Formatting-related features fail to transfer first.
- Small-model interpretability results need re-checking at scale.

#interpretability#llm#sparse-autoencoders

PAPER

relevance 9/10 · arXiv cs.AI · 2026-08-27

[arxiv.org/abs/2608.09143 ↗](#)

An 8B open-weights model ships with its eval harness

A new open-weights model in the 8B class, published together with its tokenizer, training details and the harness used to evaluate it, under a licence that permits commercial use.

- Weights, tokenizer and eval harness ship in one repository.
- The licence carries no revenue threshold.

#llm#open-weights#open-source

RELEASE

relevance 8/10 · Hugging Face Blog · 2026-08-26

[huggingface.co/blog/open-weights-8b-release ↗](#)

What “agent harness” means, and why the term keeps moving

An explainer on the scaffolding around a model in an agent system — the loop, the tool surface, the memory — and why results that vary it are hard to compare.

- Two systems on the same model can differ mostly in harness.
- Tool surface and loop shape are rarely reported.

#agents#llm#evaluation

CONCEPT

relevance 7/10 · Tavily: general · 2026-08-24

[tavily.com/blog/agent-harness-explained ↗](#)

#agents

#llm

#evaluation

#benchmarks

#tool-use

#open-source

All cards

Top tags on this page — every tag on a card is clickable.

What “agent harness” means, and why the term keeps moving

An explainer on the scaffolding around a model in an agent system — the loop, the tool surface, the memory — and why results that vary it are hard to compare.

- Two systems on the same model can differ mostly in harness.
- Tool surface and loop shape are rarely reported.

#agents

#llm

#evaluation

CONCEPT

relevance 7/10 · Tavily: general · 2026-08-24
tavily.com/blog/agent-harness-explained

Where tool use breaks down in long-horizon agent runs

An analysis of failure modes in agent runs that span many tool calls, grouping them by whether the model lost the goal, mis-read a tool result, or repeated an action it had already taken.

- Repeated actions dominate the late-run failures.
- Tool results that fail silently are the hardest to recover from.
- Run length alone predicts little without the harness.

#agents

#tool-use

#evaluation

#llm

PAPER

relevance 8/10 · arXiv cs.AI · 2026-08-22
arxiv.org/abs/2608.04471

A small harness for replaying agent benchmark runs

An open-source runner that records every tool call in an agent run and replays it against a different model, so a comparison holds the harness fixed.

- Runs are recorded as a transcript, not a score.
- Replay makes model-to-model comparison reproducible.

#agents

#benchmarks

#open-source

PROJECT

relevance 7/10 · Tavily: general · 2026-08-19
github.com/example/agent-run-replay

- #llm
- #agents
- #interpretability
- #evaluation
- #open-source
- #security
- #open-weights
- #benchmarks

Top tags on this page — every tag on a card is clickable.

AI agents rapidly exploit security vulnerabilities from patch discussions

Modern AI coding agents can now discover and exploit a vulnerability from the public discussion around its patch, well before the fix has reached most deployments.

- Patch discussions leak exploitable detail before rollout.
- The gap between disclosure and a working exploit is closing.
- Open-source projects carry the exposure first.

- #security
- #ai-agents
- #vulnerability
- #open-source
- #exploit

NEWS

relevance 9/10 · Simon Willison · 2026-08-28
[simonwillison.net/2026/Aug/28/just-a-rumour-of-a-bug/ ↗](https://simonwillison.net/2026/Aug/28/just-a-rumour-of-a-bug/)

Sparse autoencoder features transfer unevenly across model scale

A study of how features found by a sparse autoencoder on a small research model correspond to features in a deployed model of the same family, and where that correspondence breaks down.

- The clearest correspondences are the narrowest features.
- Formatting-related features fail to transfer first.
- Small-model interpretability results need re-checking at scale.

- #interpretability
- #llm
- #sparse-autoencoders

PAPER

relevance 9/10 · arXiv cs.AI · 2026-08-27
[arxiv.org/abs/2608.09143 ↗](https://arxiv.org/abs/2608.09143)

An 8B open-weights model ships with its eval harness

A new open-weights model in the 8B class, published together with its tokenizer, training details and the harness used to evaluate it, under a licence that permits commercial use.

- Weights, tokenizer and eval harness ship in one repository.
- The licence carries no revenue threshold.

- #llm
- #open-weights
- #open-source

RELEASE

relevance 8/10 · Hugging Face Blog · 2026-08-26
[huggingface.co/blog/open-weights-8b-release ↗](https://huggingface.co/blog/open-weights-8b-release)

No cards yet.

The curation run hasn't produced any cards.

#diffusion-policy

[All cards](#)

Top tags on this page — every tag on a card is clickable.

No cards tagged `#diffusion-policy` .

The whole feed was searched and nothing carries this tag.

[All cards](#)

#robotics

[All cards](#)

Top tags on this page — every tag on a card is clickable.

Searched the first 5 pages and found no cards tagged #robotics .

There are more pages behind this one. Keep going, or clear the filter.

[All cards](#)

[Next page →](#)

The feed is temporarily unavailable.

Something went wrong reading the feed.

[Try again](#)

#llm

#agents

#interpretability

#security

#open-weights

Top tags on this page — every tag on a card is clickable.

AI agents rapidly exploit security vulnerabilities from patch discussions

Modern AI coding agents can now discover and exploit a vulnerability from the public discussion around its patch, well before the fix has reached most deployments.

- Patch discussions leak exploitable detail before rollout.
- The gap between disclosure and a working exploit is closing.
- Open-source projects carry the exposure first.

#security

#ai-agents

#vulnerability

#exploit

NEWS

relevance 9/10 · Simon Willison · 2026-08-28

simonwillison.net/2026/Aug/28/just-a-rumour-of-a-bug/ ↗

An 8B open-weights model ships with its eval harness

A new open-weights model in the 8B class, published together with its tokenizer, training details and the harness used to evaluate it.

- Weights, tokenizer and eval harness ship in one repository.
- The licence carries no revenue threshold.

#llm

#open-weights

#open-source

RELEASE

relevance 8/10 · Hugging Face Blog · 2026-08-26

huggingface.co/blog/open-weights-8b-release ↗

[Next page](#) →